

Storage and Index

Instructor: Vu Tuyet Trinh



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

1

Outline

- Overview of database storage structures
- Physical database files
- Database index



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

2

2

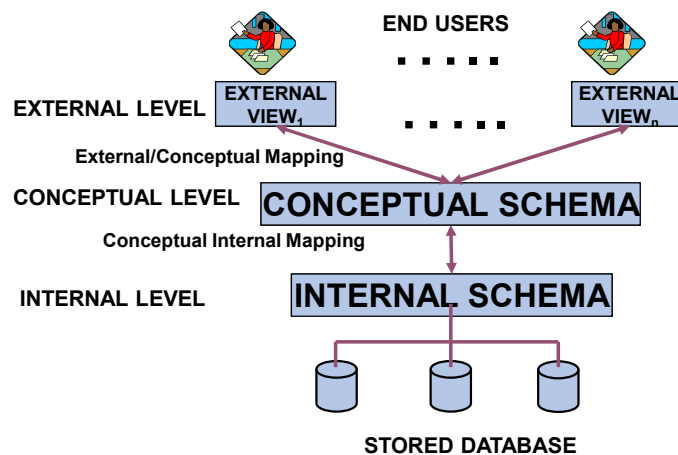
Objectives

- Upon completion of this lesson, students will be able to:
 - Understand the physical database files
 - Understand the role of database indexes

3

1. Overview of database storage structures

3-tier Schema Model (ANSI-SPARC Architecture)



4

1. Overview of database storage structures

How does Mariadb store data

```
MariaDB [(none)]> SHOW VARIABLES LIKE 'datadir';
```

| Variable_name | Value |
|---------------|-----------------|
| datadir | /var/lib/mysql/ |

```
MariaDB [student_management]> show tables;
```

| Tables_in_student_management |
|------------------------------|
| class |
| enrolled |
| faculty |
| student |

```
:/var/lib/mysql/student_management# ls -la
```

```

ql mysql 4096 Mar 12 02:05 .
ql mysql 4096 May 5 06:06
ql mysql 1547 Mar 12 02:05 class.frm
ql mysql 114688 Mar 12 02:21 class.ibd
ql mysql 65 Mar 12 01:59 db.opt
ql mysql 1466 Mar 12 02:03 enrolled.frm
ql mysql 114688 Mar 12 02:18 enrolled.ibd
ql mysql 1005 Mar 12 02:04 faculty.frm
ql mysql 98304 Mar 12 02:16 faculty.ibd
ql mysql 1101 Mar 12 02:00 student.frm
ql mysql 98304 Mar 12 02:23 student.ibd

```

the .frm table file stores the table's format
the .ibd file stores the table's data



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

5

1. Overview of database storage structures

How does Mariadb store data

- the .frm file stores the table's format

```
MariaDB [student_management]> describe student;
```

| Field | Type | Null | Key | Default | Extra |
|-------|-------------|------|-----|---------|-------|
| snum | int(11) | NO | PRI | NULL | |
| sname | varchar(40) | YES | | NULL | |
| major | varchar(30) | YES | | NULL | |
| level | varchar(10) | YES | | NULL | |
| age | int(11) | YES | | NULL | |

```
root@285e07e9458f:/var/lib/mysql/student_management# cat student.frm
?
```

```

VM?\! ?s?$??%?籐B??
??PRIMARY??InnoDB??f\P
(/?
N?

```

```
?snum?sname?major?level?age?root@285e07e9458f:/var/lib/mysql/student
```



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

6

6

1. Overview of database storage structures

How does Mariadb store data

- the .ibd file stores the table's data

```
MariaDB [student_management]> select * from student;
```

| snum | sname | major | level | age |
|------|-------------------|---------|-------|-----|
| 1 | Nguyen Van A | CS | JR | 18 |
| 2 | Nguyen Viet Cuong | History | JR | 19 |
| 3 | Nguyen Hong Ngoc | CS | JR | 19 |
| 4 | Mark Juke | History | JR | 20 |
| 5 | Elon Mulk | CS | JR | 20 |
| 6 | Donal Trump | CS | JR | 20 |
| 7 | Obama | CS | JR | 20 |
| 8 | Tan Dung | History | SR | 30 |

```
root@285e07e9458f:/var/lib/mysql/student_management# cat student.ibd
???]&!????????????????????&????????????????????&[?]Y?&??Y?&???j?&[?]??
????????????????????????i????????????????????????????????????????????
????????????????????????????????????????????????????????????????????
????????????i????????????????????????????????????????????????????????
????????????????????????????????????????????????????????j?&[?]Q????????'??E?
9infimum
supremum
.??WNguyen Van ACSJR?8?:?cNguyen Viet CuongHistoryJR! 2?@??Nguyen H
ong NgocCSJR? (0?i?iMark JukeHistoryJR? 0+?U??Elon MulkCSJR?
8-?Q?kDonal Trum
pCSJR?@'?W??ObamaCSJRH????Tan DungHistorySR?pc??Q?'??root@285e07e9458f:/var/lib/mys
```



2. Physical database files

Motivation

Magnetic disks as data storage

Primary file organizations



2.1. Motivation

- Databases typically store large amounts of data persistently on disks:
 - Databases are too large to fit entirely in main memory.
 - Disk - nonvolatile storage vs. Main memory - volatile storage
 - The cost of storage per unit is much cheaper



2.2. Magnetic disks as data storage


- A disk is a random access addressable device.
- Transfer of data between main memory and disk takes place in units of disk blocks.
- Typical disk block sizes: 4KB – 8KB.
- Disk I/O (read/write from disk to main memory) overhead is the key factor of database performance optimization.



2.2.1. Physical database design

- The process of physical database design involves choosing the **particular data organization techniques that best suit the given application requirements** (on SELECT, INSERT, UPDATE, DELETE).
- The data stored on disk is organized as files of records:
 - **Primary file organizations**: determine how the file records are physically placed on the disk, and hence how the records can be accessed.
 - **Secondary organization** or auxiliary access structure allows efficient access to file records based on alternate fields.

2.2.2. Placing File Records on Disk



© Elmasri, Ramez. *Fundamentals of database systems*. Pearson Education India, 2008

2.3. Primary file organizations

- Files of Unordered Records (Heap Files)
- Files of Ordered Records (Sorted Files)
- Hashing Techniques



2.3. Primary file organizations

- **Files of Unordered Records (Heap Files)**
 - Records are placed in the file in the order in which they are inserted
 - **INSERT:** Inserting a new record is very efficient
 - New records are inserted at the end of the file
 - **UPDATE/SELECT:** Searching for a record on any search condition is not efficient – linear search
 - **DELETE:** leaves unused space in the disk block
 - require periodic reorganization



2.3. Primary file organizations

- Files of Ordered Records (Sorted Files)
 - Physically order the records of a file on disk based on the values of one of their fields (key field)
 - **SELECT**: binary search (very fast)
 - **INSERT/DELETE/UPDATE**: more expensive



2.3. Primary file organizations

- Hash files
 - The address of the disk block in which the record is stored is the result of applying a hash function to the value of a particular field (hash field) of the record.
 - Very fast access to records for search on equality condition on the hash field.




3. Database indexes

1. What is database index?
2. Index data structures
3. B+tree
4. Spare vs. Dense index
5. Clustered vs. Non-clustered index
6. Index creation in SQL

3.1. What is database index?

- Auxiliary access structure (commonly index) allows efficient access to file records based on alternate fields




3.2. Index data structures

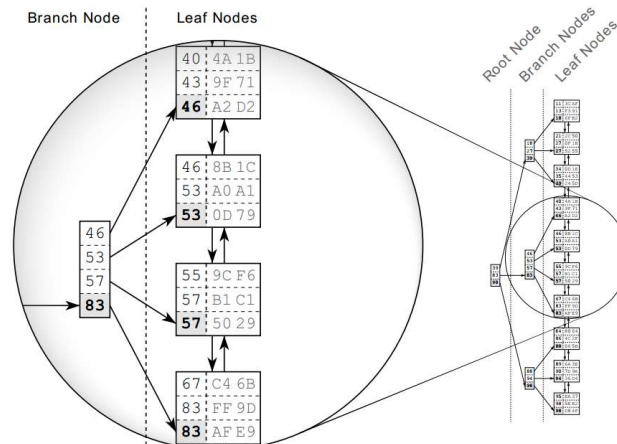
- Indexes can be implemented with different data structures.
 - B+-tree index
 - hash index
 - bitmap index (briefly)
 - dynamic hash indexes: number of buckets modified dynamically
 - R-tree: index for special data (points, lines, shapes)
 - quadtree: recursively partition a 2D plane into four quadrants
 - octree: quadtree version for three dimensional data
 - main memory indexes: T-tree, binary search tree

3.3. B+Tree

- Balanced tree of key-pointer pairs
- Keys are sorted by value
- Nodes are at least half full
- Access records for key: traverse tree from root to leaf




3.3.1. Example: B+ tree



© Gulutzan, Peter, and Trudy Pelzer. *SQL Performance Tuning*. Addison-Wesley Professional, 2003.


3.4. Spare vs. Dense index

- Sparse index
 - pointers to disk pages
 - at most one pointer per disk page
 - usually much less pointers than records
- Dense index
 - pointers to individual records
 - one key per record
 - usually more keys than sparse index optimization: store repeating keys only once, followed by pointers

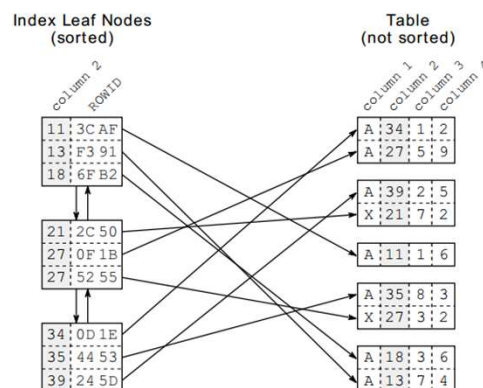


3.5. Clustered vs. Non-Clustered

- Clustered index on attribute X
 - This index controls the placement of records on disk
 - only one clustering index per table
 - dense or sparse
- Non-clustered index on attribute X
 - no constraint on table organization
 - Can have more than one index per table
 - always dense



3.5.1. Example: Non-clustered index



© Gulutzan, Peter, and Trudy Pelzer. *SQL Performance Tuning*. Addison-Wesley Professional, 2003.

3.6. Creating Index

- CREATE [UNIQUE|FULLTEXT|SPATIAL] INDEX index_name [index_type] ON tbl_name (index_col_name,...) [index_option] [algorithm_option | lock_option] ...
- index_type: USING {BTREE | HASH}



Remark

- Databases typically store data persistently on disks
 - Files of unordered records (Heap files)
 - Files of ordered records (Sorted files)
 - Hash files
- Index allows efficient access to file records based on “indexed” fields



Quiz 1.

| Quiz Number | 1 | Quiz Type | OX | Example Select |
|-------------|--|-----------|----|----------------|
| | | | | |
| Question | Does heap files support INSERT query efficiently? | | | |
| Example | A. Yes B. No | | | |
| Answer | | | | |
| Feedback | New records are appended to the end of the head file | | | |



Quiz 2.

| Quiz Number | 2 | Quiz Type | OX | Example Select |
|-------------|---|-----------|----|----------------|
| | | | | |
| Question | Are ordered files better for heavy Insert operation? | | | |
| Example | A. Yes B. No | | | |
| Answer | | | | |
| Feedback | Insertion to ordered files requires reorganizing w.r.t. new records | | | |



Summary

- Overview of database storage structures
 - 3-tier Schema Model (ANSI-SPARC Architecture)
 - How Mariadb stores data
- Physical database file structures
 - Motivation
 - Magnetic disks as data storage
 - Primary file organizations
- Database index
 - What is database indexes?
 - Index data structures
 - B+tree
 - Sparse vs. Dense index
 - Clustered vs. Non-clustered index
 - Index creation in SQL



Keywords

| | |
|--------------|--|
| Heap file | Files of Unordered Records |
| Ordered file | Physically order the records of a file on disk based on the values of one of their fields (key field) |
| Index | A data structure that improves the speed of data retrieval operations |
| B-tree | A self-balancing tree data structure that keeps data sorted |

